



Cloud Computing

DATA LAKE -

a system of data stored in natural/raw format | a single store of data including raw copies of source system data/sensor data/social data/transformed data used for reporting/visualization/advanced analytics/machine learning | includes structured data from relational databases/semi-structured data (CSV, logs, XML, JSON)/unstructured data (emails, documents, PDFs)/binary data (images, audio, video) | can be established "on premises" (within an organization's data centers) or "in the cloud" (using cloud services from Amazon/Microsoft/Google).

DATA SWAMP -

poorly managed data lakes have been facetiously called data swamps

SECURITY GROUP -

controls traffic that is allowed to reach/and leave resources | controls the inbound and outbound traffic for the instance

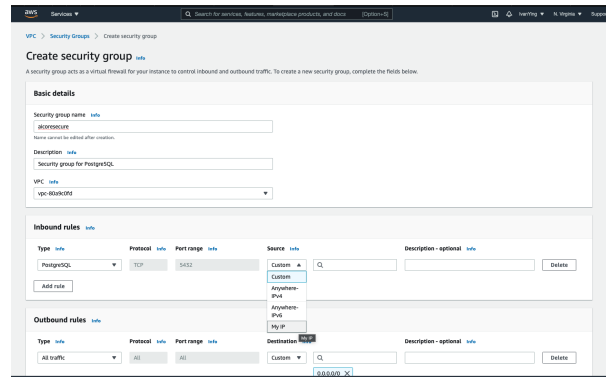
Amazon Web Services (AWS)

<https://aws.amazon.com/rds/>

Go to the console <https://us-east-1.console.aws.amazon.com/console/home?region=us-east-1#>

1. Click on the 'services' tab
2. Establish a virtual network (this will limit the range of IP addresses that can be used to access the service)
3. Click 'VCP'
4. Click 'security groups'
5. Click 'create security group'
6. Give it a name and description

1. Add an inbound rule, connect to the RDS instance from out local IP (allow the local machine to connect to the RDS ports)
2. In source, choose My IP (so you are the only one with access)
3. If you have a dynamic IP, this will work the first time, but next time your connection will be refused. In that case, change the source to Anywhere IPv4
4. Click 'create security group'



Relational Database Service (RDS)

AWS RDS is a service that allows you to create a database in the cloud. It is a highly scalable database that can be used for a variety of purposes. It allows creation of databases of different types, such as PostgreSQL, MySQL, Oracle, and Amazon Aurora

1. Click on the 'services' tab
2. Click 'RDS'
3. Click 'create database'
4. Click 'PostgreSQL'
5. Click 'Standard Create'
6. Click 'Free Tier'
7. Give it a name
8. Dont change 'Master Username'
9. Make a password
10. Check 'public access'
11. Add 'existing security group'
12. Click 'create'
13. Wait for it to be available

Take note of the Endpoint, which is the IP address of the database

```
from sqlalchemy import create_engine

DATABASE_TYPE = 'postgresql'
DBAPI = 'psycopg2'
```

```

ENDPOINT = 'datacollectionpipeline.c9nuz10uy6hn.us-east-1.rds.amazonaws.com' # Change it for your AWS endpoint
USER = 'postgres'
PASSWORD = 'Kilapila7'
PORT = 5432
DATABASE = 'postgres'

engine = create_engine(f"{DATABASE_TYPE}+{DBAPI}://{USER}:{PASSWORD}@{ENDPOINT}:{PORT}/{DATABASE}")
engine.connect()

```

Simple Storage Service (S3)

Go to the console https://us-east-1.console.aws.amazon.com/console/home?nc2=h_ct®ion=us-east-1&src=header-signin#

1. Type in 'S3'
2. Click 'create bucket'
3. Give it a name
4. Change the region
5. Type in IAM
6. Click 'users'
7. Click 'add user'
8. Make a username
9. Tick 'programmatic access'
10. Click 'next'
11. Click 'attach existing policies directly'
12. Tick 'administrator access'
13. Click 'next' until you can create a user
14. Click 'create user'
15. Download the csv (because you can only see it once)

To connect the local system to AWS

```

pip install awscli
aws configure
aws s3 ls

```

To use your AWS resources

```

pip install boto3

```

1. Install `awscli`
2. Then `aws configure`
3. Type in the info exactly as in the csv file
4. Make sure the region name matches the one on aws (eu-west-2)
5. Skip the last one by pressing enter

```
import boto3

# Uploads files to bucket
s3_client = boto3.client('s3')
# (file_name, bucket, object_name)
response = s3_client.upload_file('cat_0.jpg', 'cat-scraper', 'cat.jpg')

# Displays bucket contents
s3 = boto3.resource('s3')
my_bucket = s3.Bucket('data-collection-pipeline-bucket')
for file in my_bucket.objects.all():
    print(file.key)

# Download files from bucket
s3 = boto3.client('s3')
s3.download_file('pokemon-sprites', 'zubat/front.png', 'zubat.png')

# Download via bucket url
import requests
url = 'https://pokemon-sprites.s3.amazonaws.com/blastoise/front.png'

response = requests.get(url)
with open('blastoise.png', 'wb') as f:
    f.write(response.content)
```

Examples

```
# Link to S3 Bucket (millie-pokemon-bucket, pokemon_user)
# Upload to S3 Bucket
# !pip install boto3

import boto3

s3 = boto3.client('s3')
with open("Sprites/vulpix_front.png", "rb") as f:
    s3.upload_fileobj(f, "millie-pokemon-bucket", "vulpix_front.png")
```

6. Use `aws s3 ls` to check its working

1. Install `boto3`

2. Tell boto that you want to use a bucket

3. Upload something

4. Display contents

5. Download something

1. Click on an object in your S2 and it will give you a url

2. Use this url to download the object

Make bucket public

1. In the console, in your bucket

2. Click 'permissions'

3. Disable 'Block public access'

4. Go here

<http://awspolicygen.s3.amazonaws.com/policy>

5. Click 'type of policy'

6. Select 'S3 bucket policy'

7. Type '*' into 'principle'

8. In 'actions' select 'get object'

9. In 'Amazon Resource Name (ARN)' type 'arn:aws:s3:::your_bucket_name/*'

10. Click 'add statement'

11. Click 'Generate Policy'

12. Copy the text

13. In the console click 'permissions'

14. Edit 'bucket policy'

15. Paste the text

16. Click 'save changes'

Security Groups

SECURITY GROUP - controls traffic that is allowed to reach/and leave resources | controls the inbound and outbound traffic for the instance

When you create a VPC, it comes with a default security group.

You can create additional security groups for each VPC.

You can associate a security group only with resources in the VPC for which it is created.

For each security group, you add *rules* that control the traffic based on protocols and port numbers.

There are separate sets of rules for inbound traffic and outbound traffic.

You might set up network ACLs with rules similar to your security groups in order to add an additional layer of security to your VPC.

Elastic Compute Cloud (EC2)

“remote computers that can be employed to run code on a cluster or a single computer”

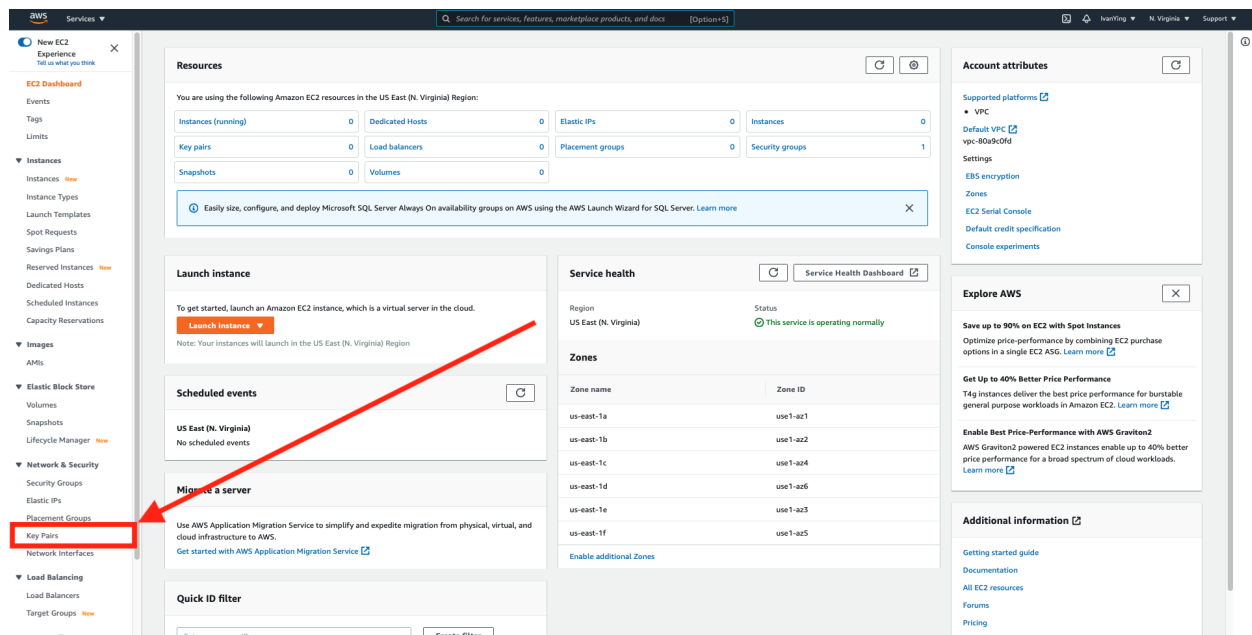
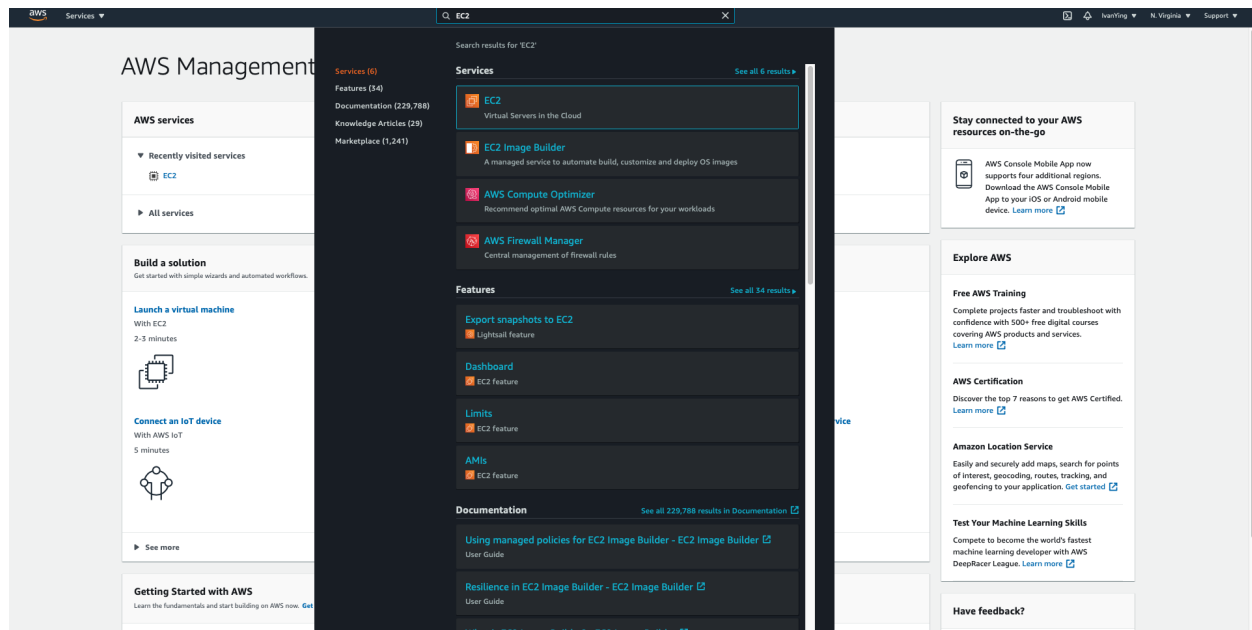
An EC2 instance can be accessed through a terminal. The instance can be accessed in the CLI by establishing an SSH connection to the instance. Once an SSH connection has been established, you can run commands as you would in your CLI.

SSH Key

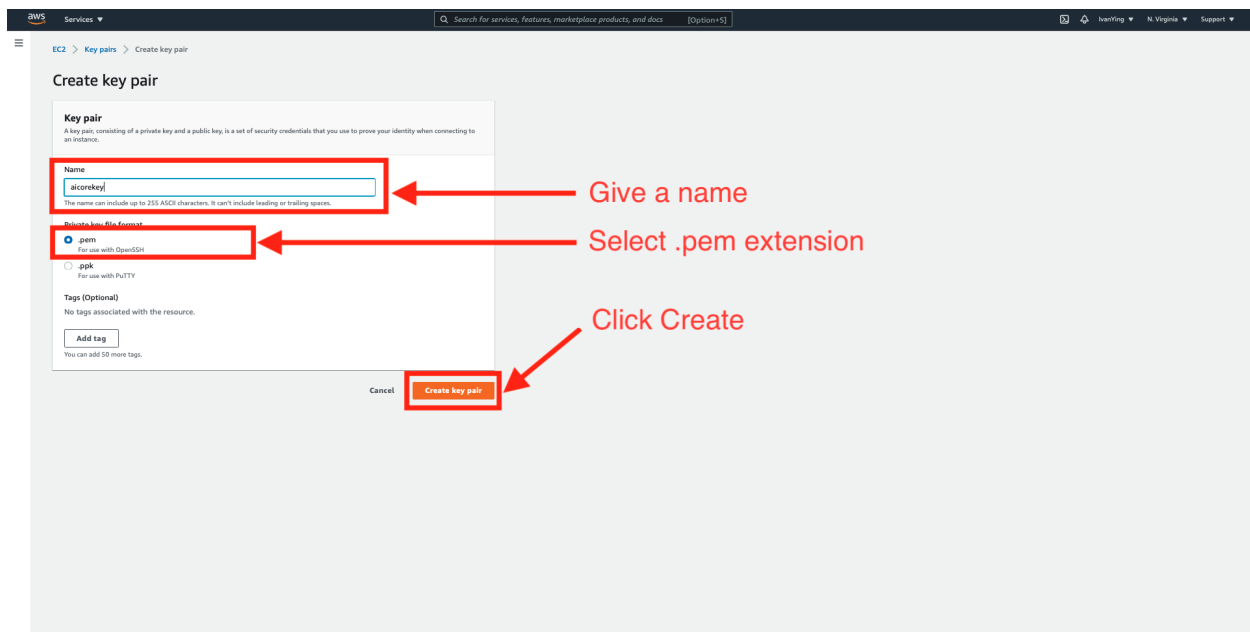
The first step is to create a key-pair file, which would enable you to connect to the instance.

To create a key pair

1. Go to the AWS console
2. Click on the EC2 section
3. Click on the **Create Key Pair** button.



4. You can provide a name, as well as an extension
5. A `.pem` file with the name of the key pair will be created. This file contains the key pair that you will use to connect to the instance, and it will be automatically downloaded immediately after you click 'Create key pair'.

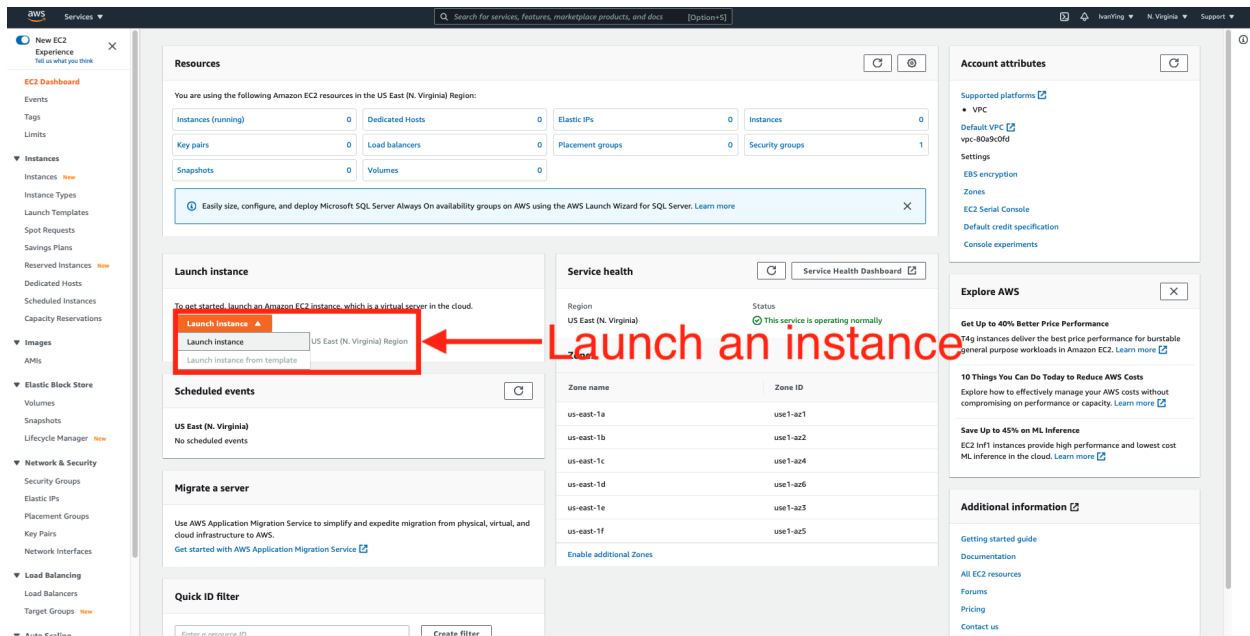


6. Ensure that the `.pem` file has only read permission. To do this, run the following command: `chmod 400 <key-pair-name>.pem`. `chmod` stands for `change mode`, and `400` is the permission. When you run the command, ensure that you are in the same directory as the `.pem` file.

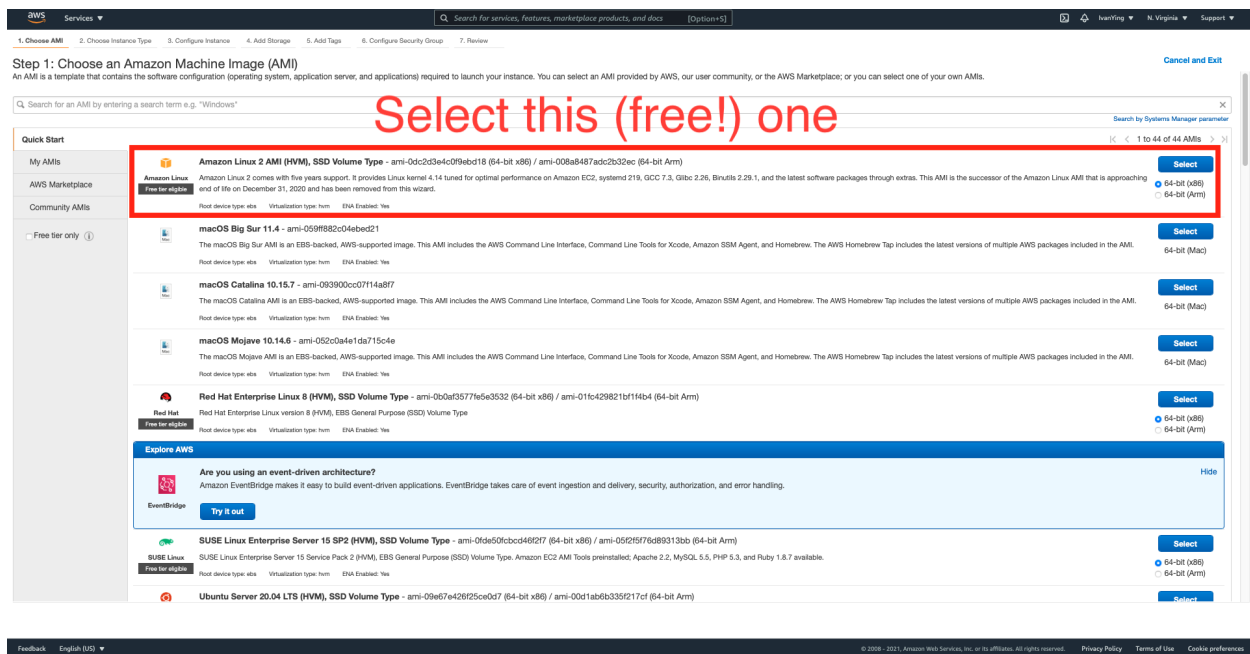
Create EC2 Instance

Remember that to enable access to the EC2 instance, a security group is required.

7. To create a security group, go to the AWS console and click on the EC2 section. You will see a list of security groups. Click on the `Create Security Group` button, and ensure that you set the correct inbound rules:
- HTTP: Anywhere IPv4
 - HTTPS: Anywhere IPv4
 - SSH: My IP
8. Once completed, remember to provide a memorable name. After creating the security group, click on the `Launch Instance` button on the EC2 main page.



9. You will be asked to choose an Amazon machine image (AMI). For this lesson, we will use the **Amazon Linux 2 AMI** image.



10. Click on Launch, and you will be prompted to select a key pair.
11. Select the key-pair file you created, and click Launch Instance. The launch process should take a few minutes.
12. Once launched, view its public DNS:

The screenshot shows the AWS Management Console interface for EC2 instances. On the left is a navigation menu with options like 'EC2 Dashboard', 'Events', 'Tags', 'Limits', 'Instances', 'Images', 'Elastic Block Store', and 'Network & Security'. The main panel displays a list of instances. The second instance, 'i-0d993512873fb107c', is highlighted and its details are shown below. In the instance list, the 'Instance state' column shows 'Running' with a green checkmark icon, which is enclosed in a red box. A red arrow points from the text 'As soon as this is running' to this box. In the instance details panel, under the 'Details' tab, the 'Public IPv4 DNS' field is highlighted with a red box. A red arrow points from the text 'You can check the DNS' to this box. The DNS value is 'ec2-3-85-43-65.compute-1.amazonaws.com'.

13. If everything works well, you should be able to connect to the instance through SSH. Ensure that your `.pem` file has read permission and that you specified the right path to the `.pem` file.
14. Connect to the EC2 instance, run the following command in the CLI: `ssh -i <key-pair-name>.pem ec2-user@<public-dns>` Remember to specify the path to the `.pem` file in the key name (unless, of course, you are in the same directory). You might be prompted to provide a fingerprint. This is the fingerprint of the key pair. You can ignore this fingerprint, and type `yes` to continue.
15. Populating the instance using secure copying, since it is empty. In your terminal, you can copy files to the EC2 instance by typing `scp -i </path/my-key-pair.pem> </path/my-file> ec2-user@<public-dns>:<path/>`

Crontab

Windows = <https://www.windowscentral.com/how-create-automated-task-using-task-scheduler-windows-10>

Crontab.guru - The cron schedule expression editor

The quick and simple editor for cron schedule expressions by Cronitor We created Cronitor because cron itself can't alert you if your jobs fail or never start. Cronitor is easy to integrate and provides you with instant alerts when things go wrong. Learn more about cron job monitoring.

 <https://crontab.guru/>

<https://vim.rtorr.com/>

Scheduled tasks can be managed using crontab. Here, we will focus on three use cases:

- `crontab -l`: List the tasks that have been scheduled.
- `crontab -r`: Remove the tasks that have been scheduled.
- `crontab -e`: Edit the tasks you want to schedule.

Before editing the crontab, we examine the structure of a crontab. First, set a schedule using five numbers:

`minutes hours day_of_month month day_of_week`

The number you give to each of these fields will set the frequency of the command. The `*` wildcard is always available for setting `every` field. Some examples:

- `* * * * *`: Every minute, every hour, every day of the month, every month, and every day of the week.
- `30 4 * * *`: Every day at 4:30.
- `* 5 * 6 *`: Every minute on all Saturdays (6), on the fifth day of the month.
- `0 0 * 11 *`: Every day at midnight in November (11).
- `0 0 1 1 *`: Yearly on the 1st of January.

The second part of the crontab comprises the command to be run. Remember that you are in the terminal. To run a python script, you should run `python3 test.py` for example.

Thus, the following command:

- `0 0 * * 5 python3 test.py`,

will run `test.py` weekly at midnight on Fridays.

Now, we add some tasks to the crontab file using the last command. Upon running it, Vim, a relatively complicated text editor, will open.

There are only two relevant commands:

- `i`: Enter Insert mode. After this, you will be able to write.
- `:wq`: Write and Quit. This saves the changes and quits the editor afterwards.

The Vim commands are provided in its console. If you are in Insert mode, simply press Esc to enter the console.

EC2 + Selenium

The created scraper might be configured to run daily or weekly.

- Problem: Running it manually might be tedious.
- Solution: Use `crontab`.
- Problem: `crontab` does not work if the computer is off.
- Solution: Run it on an EC2 instance and leave it open.
- Problem: Chrome or chromedriver for running Selenium is not available.

- Solution: Download it.
- Eventual Problem: When logging out, the script halts.
- Eventual Solution: Create a terminal multiplexor that leaves the session running.

Note: This only works with the Amazon Linux 2 AMI instance.

1. First, download Google Chrome, and determine its version. In your EC2 instance, run the following commands:

```
sudo curl https://intoli.com/install-google-chrome.sh | bash
sudo mv /usr/bin/google-chrome-stable /usr/bin/google-chrome
```

1. If no problem occurs, run the following to check the version: `google-chrome --version`

▼ For Ubuntu 20.04

1. At the time of writing, the current version is 93.x. Accordingly, we must download the corresponding chromedriver version. Since you are in a Linux instance, search for the right version on the next webpage: <https://sites.google.com/chromium.org/driver/>. Copy the link; there is no need to download the file.

```
sudo wget https://chromedriver.storage.googleapis.com/93.0.4577.15/chromedriver_linux64.zip
```

1. Unzip the downloaded file using `sudo unzip chromedriver_linux64.zip`.
2. Move that file to the PATH using `sudo mv chromedriver /usr/bin/chromedriver`.
3. You should be set to use selenium. (Remember to download `pip install selenium`).

Multiplexing

Once you start using selenium, your process will halt after the SSH is terminated. A terminal multiplexer with `tmux` presents a solution to this issue. This, in essence, will run a new terminal (or set of terminals) and keep the process running even if you disconnect from it.

To install it, run: `sudo yum install tmux`.

To improve your knowledge, experiment with `tmux`, since it has many commands. In this notebook, however, we will focus on creating terminals:

- To create a new session and connect to it, run `tmux new-session -s session_name`.
- Once inside, you can run anything as if you were in the main terminal. To detach from the session, press `Ctrl + B` and subsequently `D`.
- Create as many sessions as you want. To list the active sessions, run `tmux list-sessions``.
- To reconnect to an existing session, run `tmux attach-session -t session_name`.
- To end a session, run `tmux kill-session -t session_name`.

Once up and running, you can log out of the EC2 instance, without interrupting or halting the process.